

Predictive models

DATA1002 Lecture 10A

Alan Fekete

University of Sydney

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**). The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.


Agenda

- *Overview*
- Features
- Evaluation
- Linear Regression
- Extensions of linear regression
- K-nearest-neighbour regression
- Decision-tree classification
- “Logistic regression” classification

Use ML to train a predictive model

- Typical training

Labelled data of
many previous cases



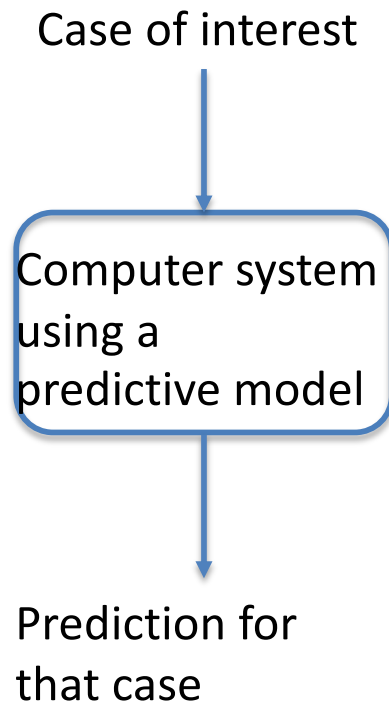
Computer system
using a
predictive model

This is *supervised* learning

- Input is a large dataset, with many previous cases
 - Eg attributes of each case
 - Also, labelled with actual value of predicted attribute
- Output of the training process is a computer system that includes a predictive model
 - Or maybe just the parameters that will be used to set up the predictive model

Deploy a predictive model

- Typical deployment



- As each case arrives, the system is run on the data about the case
 - Eg attributes of the case
- Output of the system is a prediction for some property of the case
- Eg, given input which is the contents, subject, sender etc of an email, output is either “spam” or “not spam”

More Examples

- Case is an agricultural plot growing season
 - Given data: seed-type, soil type, amount-of-rain, average-temperature, amount of nutrient1, amount of nutrient2, etc
 - Goal is to predict money-value-of-produced-crop [a *regression* task, as prediction is a quantitative attribute]
- Case is one voter
 - Given data: age, gender, job-type, salary, marital-status, number-of-dependents, etc
 - Goal is to predict party they vote for [a *classification* task, as prediction is a categorical attribute]

Agenda

- Overview
- *Features*
- Evaluation
- Linear Regression
- Extensions of linear regression
- K-nearest-neighbour regression
- Decision-tree classification
- “Logistic regression” classification

The data format

- Usually, training data is presented in “wide” logical format
 - In deployment, similar format for each case
- Each case has many attributes (also called “features”) as columns of a single row
 - In training data, the true value of the predicted attribute may be another column, or in a separate dataset that aligns with training rows
- Other terms used for the features: covariates, regressors, independent variables, input variables, control variables, explanatory variables, etc
- Other terms for the predicted feature: target, response variable, output variable, dependent variable, etc

Encoding complex data

- How do we handle data that doesn't seem like columns of integer, float, string, etc?
 - Eg data includes an image (medical scan, etc) or audio recording (eg for voice recognition)
- We may treat this as many simple-value attributes, from the basic media format
 - Eg an attribute for each separate pixel in an image,
 - Eg an attribute for sound intensity at each separate sampled moment in an audio recording

Transformed attributes

- Motivation: Many ML algorithms work best when each attribute has a similar scale, or distribution of data values
 - Eg in Australia, most adult's heights are in range 1.4 m to 2.1 m, but most annual salaries range from 15,000 to 400,000, and there are some which are much much larger
- We often transform or rescale the data of an attribute, to get a scale that works well

Transformed attributes

- Some methods that are common:
 - Transform an attribute linearly so data is mostly in range between 0 and 1, or between -1 and 1
 - Transform an attribute so distribution of data values has mean=0, variance=1 (called “z-score”)
 - Calculating the score requires the whole distribution (usually approximated using the attribute values from the training data set); the calculation can be quite expensive to perform
 - Transform data which seems to vary multiplicatively, by taking logarithm
 - Eg noise volumes expressed using dB, or earthquake intensity on Richter scale

Derived attributes

- One can include attributes that are derived from one or more of the original ones
- Eg, if original data is image, with each pixel as an attribute, we can include average-brightness-over-whole-image as a new attribute, and/or we can include an attribute is-pixel-i-on-an-edge (using edge detection calculation, based on neighbouring pixels)

Encoding a categorical attribute

- Motivation: some ML methods only work when all attributes are quantitative
 - Original data may have categorical or ordinal attributes
- “Label encoding”: each attribute value (“label”) is assigned an integer
 - Often, we assign 1 to first label we see in data, 2 to next label we see, etc
 - If attribute is ordinal, we may assign 1 to smallest label, 2 to next smallest, etc
- Alternative encoding (usually preferred in practice) is “One-hot encoding”: introduce a separate encoding attribute x_L for each original label L
 - This encoding attribute gets value 1 for item with L as its label, 0 for item with a different label
 - We say the attribute x_L is an “indicator” (or “characteristic feature”) for label L

Agenda

- Overview
- Features
- *Evaluation*
- Linear Regression
- Extensions of linear regression
- K-nearest-neighbour regression
- Decision-tree classification
- “Logistic regression” classification

Importance of evaluating a model

- There are many different ways to produce a predictive model for a given dataset
 - They may give very different models
 - We want to know whether a model is working well, or not
 - Also, we may want to choose the best from several different possibilities
 - This is a kind of hyper-parameter tuning (deciding whether to use a linear prediction or a kNN one, deciding which k to use in kNN, or how much regularisation to do, etc)
- We want to have a score that reflects the success of a model in prediction
 - This is only one aspect of the choice! We also care about computation cost, understandability, etc

Scoring a model

- There are a wide variety of different scoring definitions available
- Scikit-learn offers many of these built-in, using `sklearn.metrics` package
- Some scores are better when high (eg accuracy) and others are better when low (eg rmse). Be aware which is which
- Some scores are meaningfully compared from different algorithms on different datasets, others only can be compared between different algorithms running with exactly the same dataset
- Always: look at score on test data that was not used during training or validation
 - Otherwise scores are deceptive, and model may work much worse in practice than you expect

Scoring a regression - RMSE

- Mean-square-error is defined as square of the difference between predicted and actual value, averaged on the various items in the test dataset
 - We square each difference, so that positive and negative mis-predictions do not cancel out
 - Statistical theory that this is a good thing to optimise, for cases where errors are independent and random
- $MSE = [(predicted(test1) - actual(test1))^2 + (predicted(test2) - actual(test2))^2 + (predicted(test3) - actual(test3))^2 + \text{etc}] / (\text{number of test cases})$
- $RMSE = \text{square-root of MSE}$
 - This has same units as target attribute
- But hard to interpret, except by comparing two scores on the same test set
 - We desire a small value for RMSE

Scoring a regression - R^2

- Another score often used is called R-squared
- R^2 is at most 1
 - $R^2 = 1$ means excellent prediction
 - Despite its name, R^2 can be negative!
 - Predictions in social science rarely have R^2 above 0.5, but predictions in physical sciences often can be much closer to 1 (eg 0.9)
- See https://en.wikipedia.org/wiki/Coefficient_of_determination

Scoring a classifier – accuracy

- Accuracy of a classifier is defined as the fraction of test items where the prediction is correct (ie how many i where $\text{predict}(\text{test}_i) = \text{actual}(\text{test}_i)$)
- Warning: when the data is unbalanced (many more examples in one class compared to others, eg most patients do not have a disease), this is not very useful
 - high accuracy can come by just always predicting most common case (ignoring the data attributes)

Confusion matrix

- Count how many test cases there are, for each combination of predicted-label and actual-label
 - Record counts in a matrix
- Eg if we predict cat, and animal is really dog, this is counted in matrix entry at intersection of row predict=cat and column where actual=dog
- $\text{Accuracy} = (\text{sum of diagonal entries}) / (\text{sum of all entries})$
- Full matrix helps us see what kinds of mistake our classifier is making

Scoring a binary classifier

- For a binary classifier (only two labels, eg item has-property-of-interest or item doesn't-have-property-of-interest)
- Name confusion matrix entries as

Predicted \ Actual	actually has property	Actually doesn't have property
Positive (predicted to have property)	TP	FP (number of False Positive cases)
Negative (predicted as not-with-property)	FN (number of False Negative cases)	TN

Scores for binary classifier

- Many different names used in different fields of scholarship, for various ratios
- Sensitivity, recall, hit-rate: $TP/(TP+FN)$
- Precision: $TP/(TP+FP)$
- Selectivity: $TN/(TN+FP)$
- $F1 = 2*TP/(2*TP+FN+FP)$
- In real applications, different situations may have very different impact
 - Eg predict-disease when no-actual-disease causes some inconvenience of extra tests, but predict-no-disease when actually has-disease may lead to patient death, and predict-disease when has-disease may save the life
 - So we want to choose predictive model that makes impact as good as possible
 - Choose to optimise the score that suits the domain characteristic

Agenda

- Overview
- Features
- Evaluation
- *Linear Regression*
- Extensions of linear regression
- K-nearest-neighbour regression
- Decision-tree classification
- “Logistic regression” classification

Linear predictive models

- Prediction = $w_1 * \text{case}.x_1 + w_2 * \text{case}.x_2 + w_3 * \text{case}.x_3 + \text{etc}$
 - Here x_1, x_2 etc are attributes of a case
 - w_1, w_2 etc are “weights” that determine which predictive model is being used, within this family
 - Other notation: weights are also called coefficients, and are often denoted as β_1, β_2 etc
- Each attribute contributes linearly to the prediction
 - A given increase in the attribute, leads to same change in predicted value, no matter what value the attribute has, nor what values other attributes have

Limitations

- All attributes must be quantitative
 - Use encoding to convert any attributes that are originally categorical or ordered
- Choice of model is most influenced by the attributes with greatest variation
 - It is wise to transform and rescale each attribute, so they are comparable
- Before training a linear model, look at data to be convinced it seems linear
 - If not, maybe transform some attributes to adjust
- Linear models are quite fragile if data contains errors or outliers
 - It's wise to remove these before training a model

Which model in the family?

- The simplest (and traditional) choice is to find the set of weights for which the mean-square-error (MSE) is smallest *on the training set*
 - That is, minimize $[(\text{predicted}(\text{train1}) - \text{actual}(\text{train1}))^2 + (\text{predicted}(\text{train2}) - \text{actual}(\text{train2}))^2 + (\text{predicted}(\text{train3}) - \text{actual}(\text{train3}))^2 + \text{etc}] / (\text{number of training items})$
 - Here each training instance train_i is itself a vector of attribute values
- There is a direct calculation for this using matrix operations
- But heuristic search is often used when there are large amounts of data to train on
 - Eg Stochastic gradient descent (“SGD”)

Scikit-learn training and deployment

```
import pandas as pd
from sklearn import linear_model
from sklearn import metrics
from sklearn.model_selection import train_test_split

df = pd.read_csv('CAvideos_stat.csv')
X = df.values[:, 1:3]      # slice DataFrame for input variables
y = df.values[:, 3]        # slice DataFrame for target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
random_state=42)

# train linear regression predictor to fit training data
regr = linear_model.LinearRegression().fit(X_train, y_train)

# The coefficients
print('Coefficients:')
print(regr.coef_)

# Use the model to predict y from X_test
y_pred = regr.predict(X_test)
```

This code example

(from Grok ML module)

uses data stored as pandas dataframe

Scikit-learn evaluation

```
import pandas as pd
from math import sqrt
from sklearn import metrics

# Somehow produce a model regr
# Use the model to predict y from X_test
y_pred = regr.predict(X_test)

# y_test is dataframe of actual values for X_test
# mean squared error
mse = metrics.mean_squared_error(y_test, y_pred)
print('Root mean squared error (RMSE):', sqrt(mse))
# R-squared score
print('R-squared score:', metrics.r2_score(y_test, y_pred))
```





Interpretation of weights

- Coefficient of x_1 gives “amount two cases are predicted to differ in target-attribute, if they are the same on all other attributes of the dataset, and differ by 1 in value for x_1 ”
- Do not treat this as causal without much more evidence
 - Do not say “if we increase x_1 by 1, then target attribute will increase by w_1 ”
 - one often can’t simply change one attribute while keeping others the same
 - Eg suppose we predict salary from height!
 - Decision about which attributes are included in the model, can even alter the sign of the coefficient

Warnings

- Remember how misleading it is, to look at score on data that is used during training (or while validating to tune hyper-parameters)
- Don't extrapolate with the model to cases that are quite unlike training data
 - Instead, only predict for data that is similar to training cases
 - Eg if training data is mostly in developed countries, don't expect model to hold for developing world
 - Eg if training data is crop yields in tropics, don't expect model to hold in temperate regions
- The software makes it very easy to produce a model, but the model may predict quite badly (eg if real world is not close to linear, if training data has outliers, etc)

Anscombe's quartet (all have same best-fit linear predictor)



Agenda

- Overview
- Features
- Evaluation
- Linear Regression
- *Extensions of linear regression*
- K-nearest-neighbour regression
- Decision-tree classification
- “Logistic regression” classification

Polynomial regression

- If we expect target attribute to depend non-linearly on attribute x_i , a common technique is to include x_i , and then also x_i^2 , x_i^3 etc as extra (derived) attributes in a linear regression
 - Up to some maximum power of x_i
- Warning: only do regression with powers of x_i , when you have good reason to expect a polynomial impact of the degree you are including
 - It is very easy to overfit training data with polynomial model

Interactions between features

- If we expect an interaction between x_1 and x_3 , we might include $x_1 * x_3$ as an extra derived attribute
 - This works best if one of the attributes is indicator variable (value is 0 or 1, or sometimes -1/1 is used)
- Warning: the number of interactions one might decide to include grows rapidly with number of attributes

Regularisation

- With many attributes in a dataset, overfitting becomes quite likely
- The search for a good model can be shaped to reduce this danger, by penalizing models where many weights are large
- Many ways to do this, often controlled by hyper-parameters
- Eg in scikit-learn

```
from sklearn import linear_model  
regr = linear_model.Ridge(alpha=.5).fit(X_train, y_train)
```

Agenda

- Overview
- Features
- Evaluation
- Linear Regression
- Extensions of linear regression
- *K-nearest-neighbour regression*
- Decision-tree classification
- “Logistic regression” classification

Overview

- For a given case where we want a prediction, we find the k nearest members of the training set
 - Measure “distance” between two rows is based on the difference in each attribute, combined somehow
 - Eg “Euclidean distance” between case and traini is $\sqrt{(\text{case}.x_1 - \text{traini}.x_1)^2 + (\text{case}.x_2 - \text{traini}.x_2)^2 + \dots}$
- Prediction for case is *weighted average* of the actual target value on each of the nearest training instances, combined with weight which is $1/(\text{distance of that training instance from the case})$
- k (the number of neighbours used) is a hyper-parameter of this regression technique which is often called “kNN” regression

Some strengths and limitations

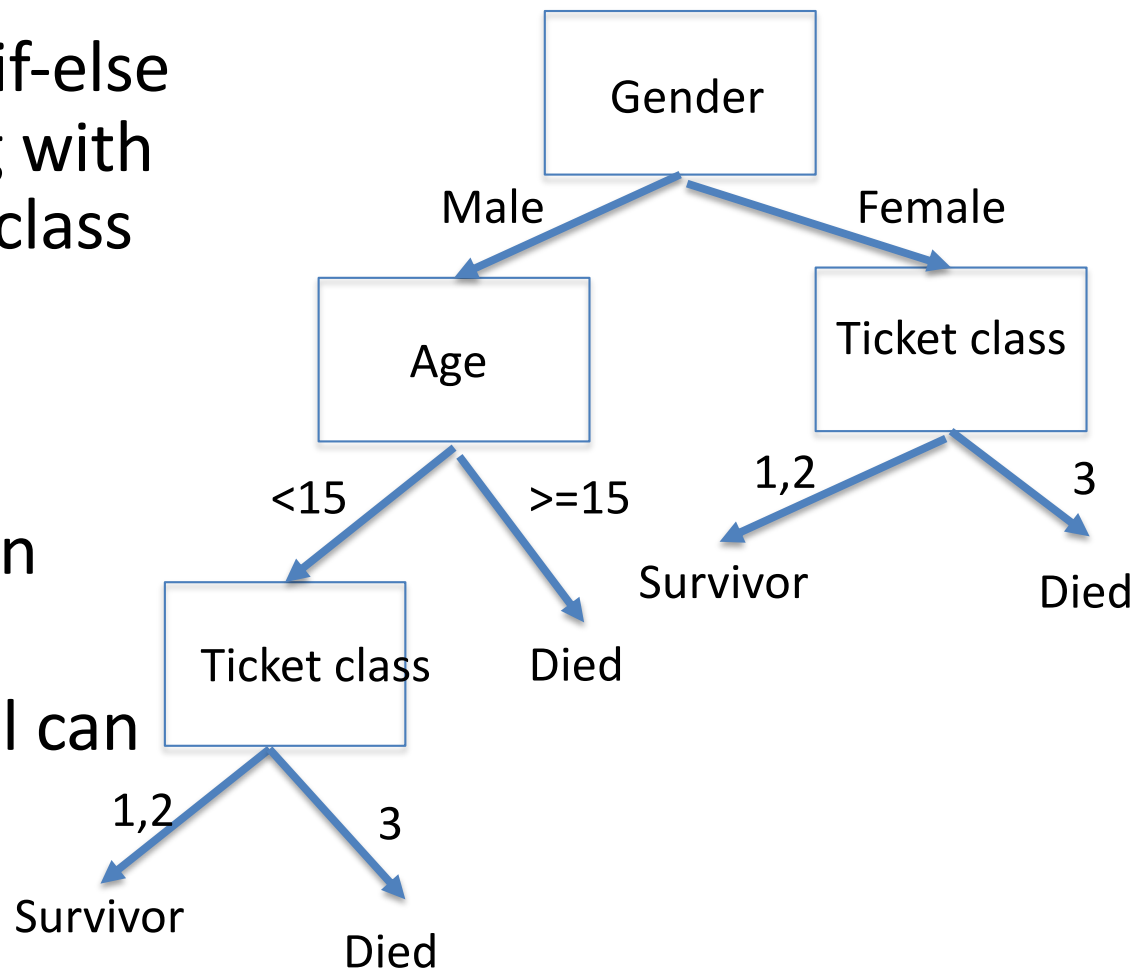
- kNN works well with data that has quite complicated patterns (not just nearly linear)
- kNN is not disturbed much by a few outliers/errors in data
- Deploying a kNN predictive model takes a lot of computation and storage

Agenda

- Overview
- Features
- Evaluation
- Linear Regression
- Extensions of linear regression
- K-nearest-neighbour regression
- *Decision-tree classification*
- “Logistic regression” classification

Overview of the family of models

- Essentially, a deeply nested collection of if-else logic choices, ending with the assignment of a class label
 - Like the game “20 questions”
- Each choice condition can be different
- The predictive model can be easy for users to understand



Possible classifier for Titanic passenger survival

Variety of models

- Each decision looks at the value of one attribute, and chooses next step based on comparing the value to model parameters
 - For split on a quantitative attribute, partition the range of possible values
 - For split on categorical attribute, each possible value can be a separate choice

Scikit-learn code

```
import pandas as pd
from sklearn import metrics
from sklearn import tree
from sklearn.model_selection import train_test_split

df = pd.read_csv('dataR2.csv')
[num_row, num_var] = df.shape
X = df.values[:, 0:num_var - 1] # slice dataFrame for input variables
y = df.values[:, num_var - 1] # slice dataFrame for target variable
X_train, X_test, y_train, y_test = \
    train_test_split(X, y, test_size=0.3, random_state=42)
clf = tree.DecisionTreeClassifier().fit(X_train, y_train)

y_pred = clf.predict(X_test)

print('Calculate the accuracy using the test data')
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Choosing the best model

- CART algorithm:
- Find the choice of one splitting attribute and thresholds that produces the “lowest impurity” division
 - The subsets defined by this split are as close to uniform-label as possible (and different from one another)
- Then treat each subset as a new dataset and find its best split, etc

Warnings

- Decision tree classifiers are quite prone to overfitting
 - So we often restrict chosen tree so it is not too deep, etc
- When there are many attributes, the search for best decision tree can be very slow

Agenda

- Overview
- Features
- Evaluation
- Linear Regression
- Extensions of linear regression
- K-nearest-neighbour regression
- Decision-tree classification
- *“Logistic regression” classification*

Overview

- A logistic regression predicts a probability value “how likely is it that a case with these attributes, is in the class”?
- To use this for binary classification, we assign the case to the class, when the probability predicted is greater than 0.5
 - If there are multiple possible classes, predict a probability for each, and then assign case to whichever class has highest probability

Further learning

- Many wikipedia articles, start from:
https://en.wikipedia.org/wiki/Supervised_learning
- Data Science from Scratch (2nd ed), by J. Grus, O'Reilly 2019
- Hands-On Machine Learning with Scikit-Learn & TensorFlow, by A. Géron, O'Reilly 2017
- The Master Algorithm, by P. Domingos, Penguin 2015
- USyd unit COMP3308 Introduction to Artificial Intelligence
- For scikit-learn, see <https://scikit-learn.org/>